

Evaluation Plan

Discovering Your Narrative Voice — Measuring Whether the Revision Actually Worked

The Content Quality Rubric judges whether content meets a standard before it publishes. This plan is different: it defines how to find out whether the revised module actually performs better once real learners are using it — not whether it should work, but whether it does.

1. Framework

Structured around Kirkpatrick's four levels of training evaluation, adapted for this platform's constraints. Each level gets harder to measure as you go down the list — that's expected, and this plan says so explicitly rather than pretending every level is equally reachable with the data actually available.

Level	Question It Answers	Measurable Here?
1. Reaction	Did learners find the module usable and tolerable to get through?	Yes — directly, via existing feedback screen and platform analytics.
2. Learning	Did learners actually build the target skill (narrative voice, character construction)?	Limited — activities are written on paper by design (see Content Outline, Section 5), so work product isn't visible to evaluate directly without undermining the privacy the design intentionally protects.
3. Behavior	Does the skill transfer to real writing outside the module (letters, reentry documents)?	Not measurable from platform data alone — would require learner self-report, out of scope for this module in isolation.
4. Results	Does this move any broader outcome (literacy gains, reentry readiness)?	Out of scope for a single module — belongs at a program or curriculum level, not this evaluation plan.

2. What This Plan Actually Measures

Given the Level 2-4 limits above, this plan focuses on what's honestly available: Level 1 reaction data, plus platform-level behavioral signals that indirectly test whether the specific fixes from the Content Review Case Study worked as intended.

Behavioral signals (platform analytics, aggregate only)

Signal	What It Tests
Completion rate for the full module	Baseline — did the platform-independence and video-gate fixes reduce the number of learners who start but never finish?
Drop-off point, if any (which screen a learner exits on and doesn't return)	Directly tests the Task & Content Analysis hypothesis: Screens 5 and 8 were the highest-risk failure points in the original. If drop-off there decreases post-revision, that hypothesis is confirmed.
Skip-button usage rate on Screens 2, 7, and 10	Tests whether the learner-control fix mattered at all. High skip usage means learners value the choice; near-zero usage would suggest the

Signal	What It Tests
	original lock wasn't the barrier the audit assumed — worth knowing either way.
Time spent on Screens 3-4 (accordion) relative to module average	Indirect check on whether the chunking approach is genuinely reducing load or just relocating it — unusually long dwell time could indicate confusion rather than engagement.

Self-reported reaction (Screen 13 feedback field)

The module already includes an open-text "How would you improve this lesson?" field. This plan treats that field as a real data source, not just a formality — responses should be reviewed as a set after each cohort, coded for recurring themes, and fed back into the next revision cycle rather than collected and left unread.

3. Evaluation Constraints — This Measurement Has to Follow Its Own Rules

An evaluation plan built for this population can't quietly reintroduce the surveillance and comparison problems the revised module was built to remove. These constraints apply to the measurement itself, not just the content being measured.

- All behavioral data reviewed in aggregate only. No individual learner's completion time, skip usage, or drop-off point should be visible to facility staff at a level that could read as monitoring compliance or effort.
- No data from this evaluation should be usable to penalize a learner — not for slow completion, not for skipping a video, not for an incomplete run. If a data point could plausibly be used punitively, it doesn't belong in a report that leaves the instructional design team.
- Feedback field responses (Screen 13) should be reviewed anonymized and in aggregate, not attributed to individual learners in any report.

4. Cadence

Checkpoint	What Happens
After first cohort (approx. 50-100 completions)	Pull completion rate, drop-off points, and skip usage. Compare against pre-revision baseline if available. Read the first batch of feedback responses.
Quarterly thereafter	Re-check the same metrics for drift. A module that performed well at launch can degrade in perceived relevance over time even with no content changes.
Any time the module is revised again	Re-baseline. Don't compare post-revision data against a stale pre-revision number from a different platform version.

5. What Would Trigger a Re-Audit

- Drop-off at Screens 5 or 8 doesn't meaningfully improve post-revision — would suggest the embedded-content fix didn't address the actual barrier, and the Task & Content Analysis hypothesis needs revisiting.

- Skip usage on video screens is near-zero — would suggest either the video content itself was already well-received (a good sign) or the skip option isn't visible/discoverable enough (a design problem worth a follow-up accessibility check).
- Feedback responses cluster around a theme not addressed anywhere in this document set — evidence of a gap the original audit didn't catch, which the Content Quality Rubric should be updated to include.